


# Outlines

Speech2Text  
Whisper-OpenAI



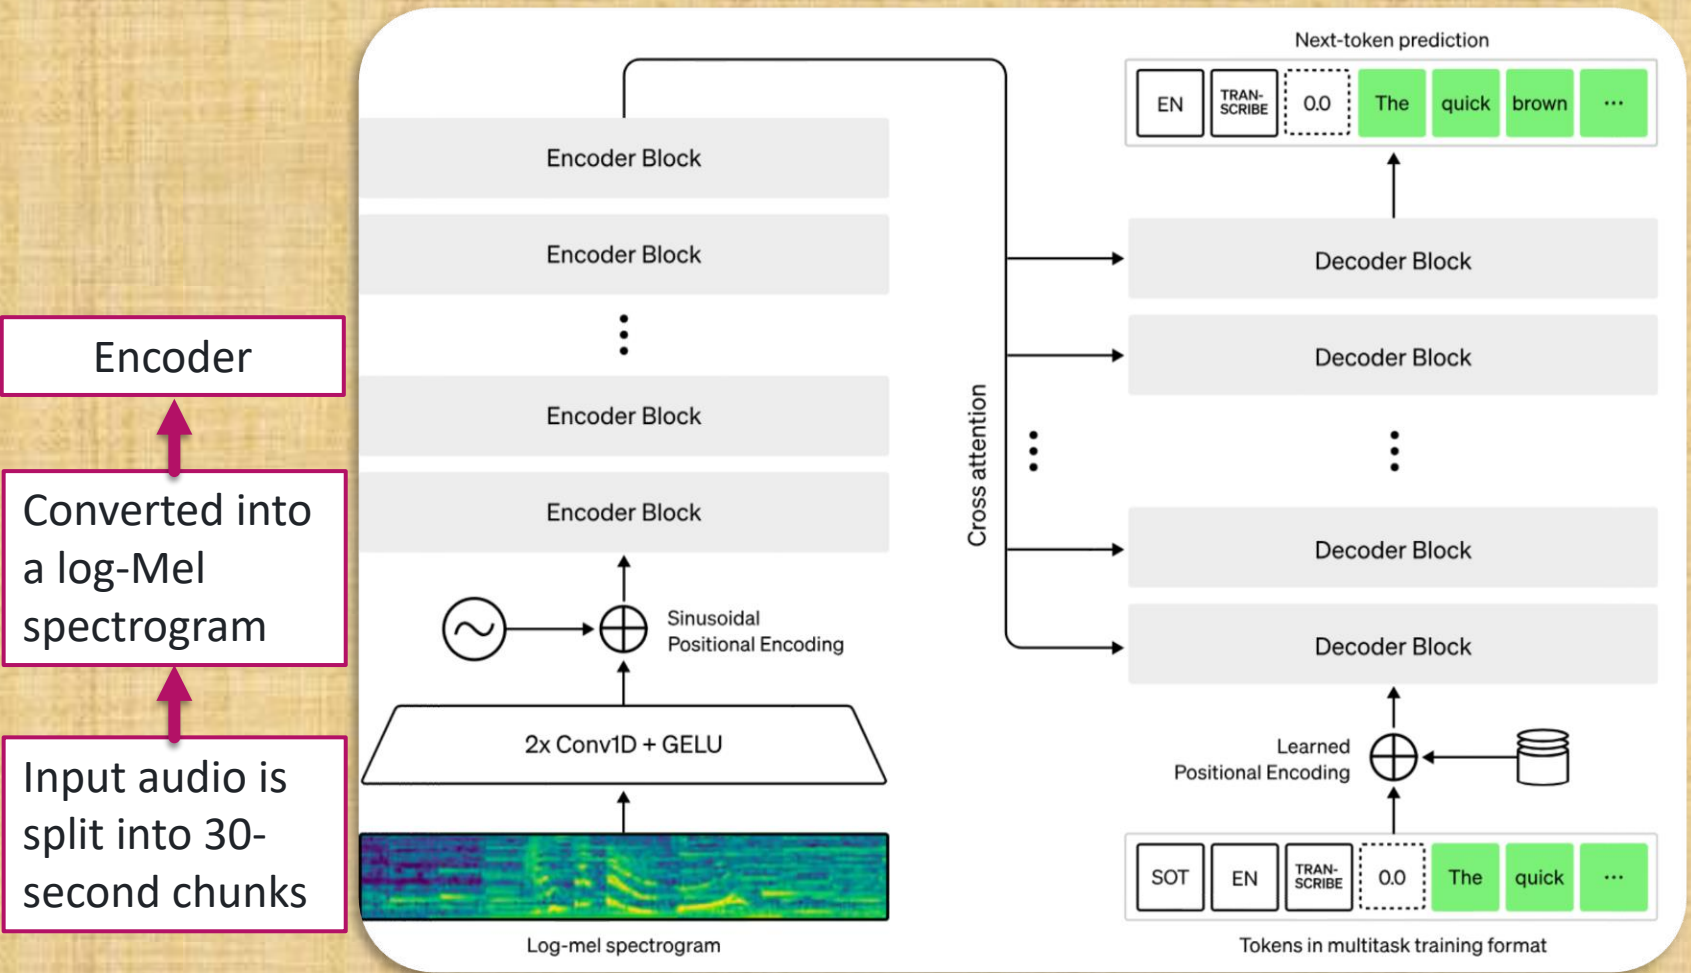
NVIDIA NeMo  
Framework



Quantization  
TF32, FP16



Speech2Text  
Whisper  
OpenAI



Multilingual Speech Recognition  
to-English Speech Translation  
Language Identification  
phrase-level timestamps  
Voice Activity Detection

Python 3.9.9  
PyTorch 1.10.1  
OpenAI's tiktoken  
Ex. Py 3.8-3.11

```
pip install -U openai-whisper  
sudo apt update  
sudo apt install ffmpeg  
pip install setuptools-rust
```

# Speech2Text Whisper OpenAI

## Whisper

<https://openai.com/research/whisper>

<https://github.com/openai/whisper/tree/main>



## My Github with Dockerfile and example

[https://github.com/HFarkhari/Whisper\\_OpenAI](https://github.com/HFarkhari/Whisper_OpenAI)

## Docker Image (Large-v2 model included)

[https://hub.docker.com/r/hfarkhari/whisper\\_openai\\_speech2text](https://hub.docker.com/r/hfarkhari/whisper_openai_speech2text)



```
docker pull hfarkhari/whisper_openai_speech2text:large-v2
```

```
docker run --gpus all -p 8888:8888 -it --rm hfarkhari/whisper_openai_speech2text:large-v2
```

<http://127.0.0.1:8888>      password: 123

NVIDIA NeMo  
Framework

Speech2Text  
Whisper  
OpenAI

<https://huggingface.co/openai/whisper-large-v2>



HUGGING FACE

Size	Parameters	English-only model	Multilingual model	Required VRAM	Relative speed
tiny	39 M	<code>tiny.en</code>	<code>tiny</code>	~1 GB	~32x
base	74 M	<code>base.en</code>	<code>base</code>	~1 GB	~16x
small	244 M	<code>small.en</code>	<code>small</code>	~2 GB	~6x
medium	769 M	<code>medium.en</code>	<code>medium</code>	~5 GB	~2x
large	1550 M	N/A	<code>large</code>	~10 GB	1x

Large V2 ~10GB VRAM (GPU)

More Info

<https://www.assemblyai.com/blog/how-to-run-openais-whisper-speech-recognition-model/>



## Tacotron2 - Text 2 Speech (TTS)

<https://paperswithcode.com/task/text-to-speech-synthesis>

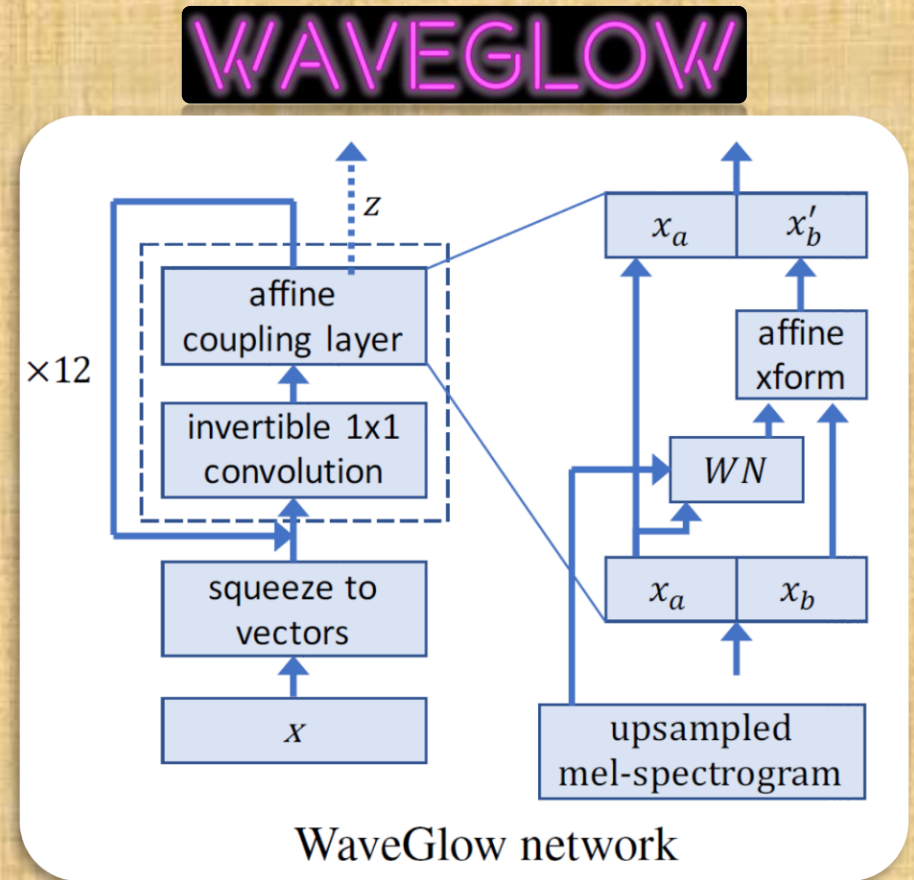
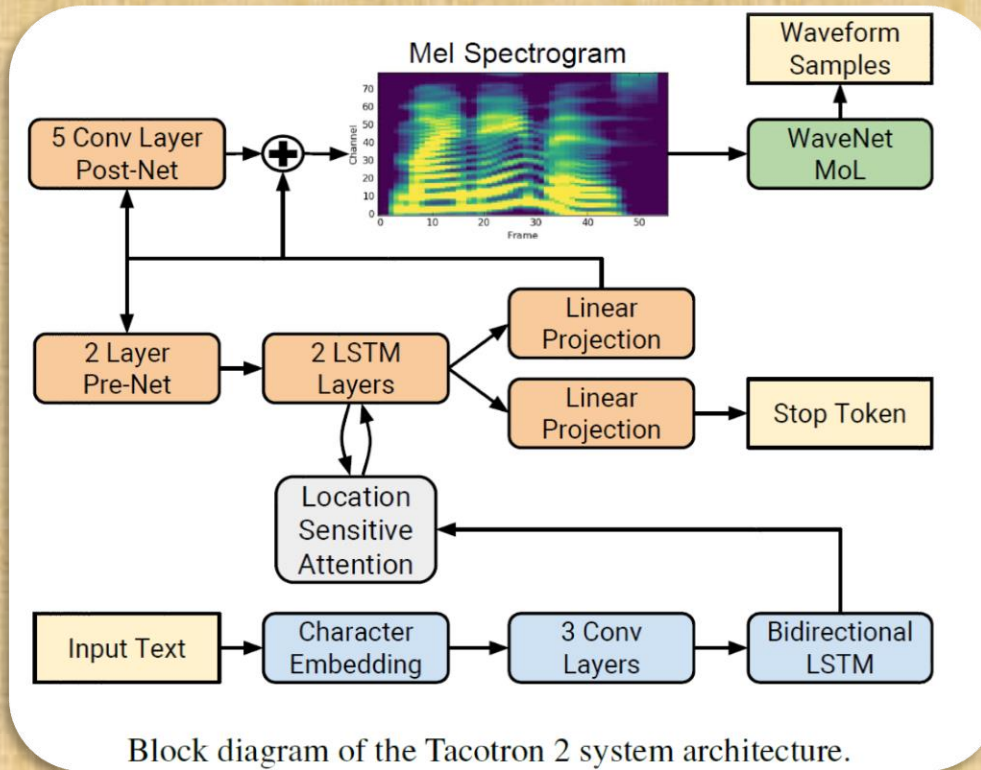
<https://github.com/NVIDIA/tacotron2>

[https://pytorch.org/hub/nvidia\\_deeplearningexamples\\_tacotron2/](https://pytorch.org/hub/nvidia_deeplearningexamples_tacotron2/)

## WaveGlow: a Flow-based Generative Network for Speech Synthesis

<https://github.com/NVIDIA/waveglow>

<https://nv-adlr.github.io/WaveGlow>



# NVIDIA NeMo Framework

## NVIDIA NeMo Framework

<https://github.com/NVIDIA/NeMo/tree/main/examples>



- Automatic Speech Recognition (ASR)
- NLP (Language Modeling, Information Retrieval, Machine Translation, Question Answering, Text Classification, ...)
- Speech Intent Classification and Slot Filling on SLURP Dataset
- TTS (FastPitch, Tacotron2, WaveGlow, Spectrogram Enhancer, ... )

<https://catalog.ngc.nvidia.com/orgs/nvidia/containers/nemo>

```
docker pull nvcr.io/nvidia/nemo:23.04
```

```
docker run --gpus all -it --rm -p 8888:8888 nvcr.io/nvidia/nemo:23.04
```

```
docker run --gpus all -it --rm --shm-size=8g -p 8888:8888 -p 6006:6006 --ulimit memlock=-1  
--ulimit stack=67108864 nvcr.io/nvidia/pytorch:23.04-py3
```



Quantization  
TF32, FP16

# NVIDIA NeMo Framework

Speech2Text  
Whisper  
OpenAI

## NVIDIA NeMo Framework

<https://github.com/NVIDIA/NeMo/tree/main/examples>



## Tutorials

<https://docs.nvidia.com/deeplearning/nemo/user-guide/docs/en/stable/starthere/tutorials.html>

- Automatic Speech Recognition (ASR)
- NLP (Language Modeling, Information Retrieval, Machine Translation, Question Answering, Text Classification, ...)
- Speech Intent Classification and Slot Filling on SLURP Dataset
- TTS (FastPitch, Tacotron2, WaveGlow, Spectrogram Enhancer, ... )

<https://catalog.ngc.nvidia.com/orgs/nvidia/containers/nemo>

```
docker pull nvcr.io/nvidia/nemo:23.04
```

```
docker run --gpus all -it --rm -p 8888:8888 nvcr.io/nvidia/nemo:23.04
```

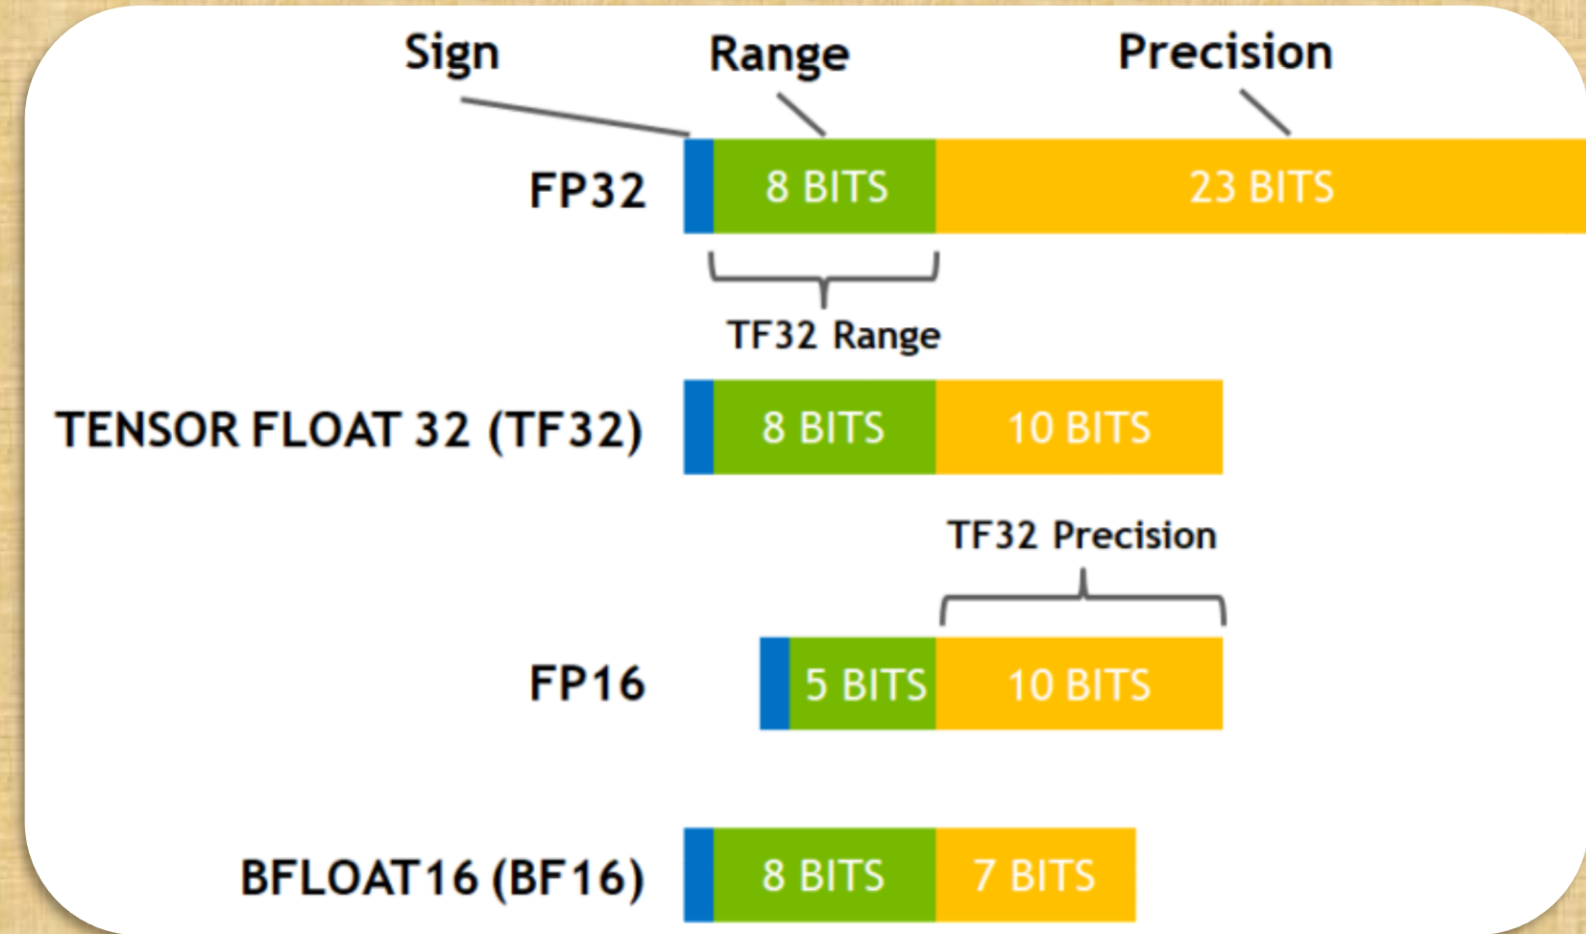
```
docker run --runtime=nvidia -it --rm -v --shm-size=16g -p 8888:8888 -p 6006:6006  
--ulimit memlock=-1 --ulimit stack=67108864 nvcr.io/nvidia/nemo:23.04
```





# FP32, TF32, FP16, BF16

<https://developer.nvidia.com/blog/accelerating-ai-training-with-tf32-tensor-cores/>



Quantization  
TF32, FP16

NVIDIA NeMo  
Framework



# FP32, TF32, FP16, BF16

<https://developer.nvidia.com/blog/accelerating-ai-training-with-tf32-tensor-cores/>

## TensorFloat32 (TF32) support

- Accelerating FP32 **Only convolutions** and **matrix multiplications**.
- All storage in memory and other operations remain completely in **FP32**, only **convolutions** and **matrix-multiplications** convert their inputs to **TF32** right before multiplication.
- TF32 mode is the default option for AI training with 32-bit variables on **Ampere GPU architecture (A100)**.
- Rounds FP32 inputs to TF32, computes the products **without loss of precision**, then accumulates those products into an FP32 output.
- **TF32** is only exposed as a Tensor Core operation mode, **not a type**.
- achieves the **same accuracy** as FP32 training, requires **no changes** to **hyperparameters** for **training scripts**

Quantization  
TF32, FP16


# FP32, TF32, FP16, BF16

<https://developer.nvidia.com/blog/accelerating-ai-training-with-tf32-tensor-cores/>

## TensorFloat32 (TF32) support

- **FP16** gives a further **speedup of up to ~2x**, as 16-bit Tensor Cores are 2x faster than **TF32 mode**

Quantization  
TF32, FP16



NVIDIA NeMo  
Framework

## TensorFloat32 (TF32) Tips

- **TF32** is the **default mode** for AI on A100 when using the NVIDIA optimized deep learning framework containers for TensorFlow, PyTorch, and MX Net, starting with the **20.06**.
- PyTorch 1.7, TensorFlow 2.4, nightly builds for MXNet 1.8 (minimum version).
- cuDNN version 8.0 and greater.
- cuBLAS version 11.0 and greater.
- cuSOLVER
- cuTENSOR version 1.1.0 and greater.
- **BF16** is introduced as Tensor Core math mode in **cuBLAS 11.0** and as a numerical type in **CUDA 11.0**.
- Deep learning frameworks and AMP will support **BF16** soon.



### TensorFloat32 (TF32) Tips

- TF32 mode accelerates **convolution** and **matrix-multiply layers**, including **linear** and **fully connected** layers, **recurrent cells**, and **attention blocks**.
- TF32 does not accelerate layers that operate on **non-FP32** tensors, such as **16-bits**, **FP64**, or **integer** precisions.
- TF32 also does not apply to layers that are not **convolution** or **matrix-multiply** operations (for example, **batch normalization**), as well as **optimizer** or **solver** operations.
- Tensor storage is not changed when training with TF32. Everything remains in FP32, or whichever format is specified in the script.

## TensorFloat32 (TF32) Tips

### PERFORMANCE

How does one know tensor cores were used?

```
import torch
import torch.nn
```

```
bsz, inf, outf = 256, 1024, 2048
```

```
tensor = torch.randn(bsz, inf).cuda().half()
layer = torch.nn.Linear(inf, outf).cuda().half()
layer(tensor)
```

FP16 input  
FP16 output

Running with:

```
nvprof python test.py
```

Produces (among with other output):

```
37.024us  1  37.024us  37.024us  37.024us  volta_fp16_s884gemm_fp16.
```

Tensor core 884



<https://youtu.be/tAIakfEt-tI>

Quantization  
TF32, FP16

# TensorFloat32 (TF32) Tips

```
nvprof --log-file results/nvprof_log.txt python train.py
```

Avg	Min	Max	Name
1.898us	27.456us	3.7446ms	volta_fp16_s884gemm_fp16_128x64_ldg8_f2f_nn
3.74us	58.335us	6.2330ms	volta_fp16_s884gemm_fp16_256x128_ldg8_f2f_nn
1619ms	152.83us	5.5497ms	volta_fp16_s884gemm_fp16_256x128_ldg8_f2f_nt
1.079us	41.760us	1.0913ms	volta_fp16_s884gemm_fp16_256x64_ldg8_f2f_nn
1.27us	36.767us	1.4038ms	volta_fp16_s884gemm_fp16_128x128_ldg8_f2f_nt
3465ms	42.560us	4.1843ms	volta_fp16_sgemm_fp16_128x64_nt

Quantization  
TF32, FP16



# TensorCore Performance Guidance

- **Requirements to trigger TensorCore operations:**
  - Convolutions:
    - Number of input channels a multiple of 8
    - Number of output channels a multiple of 8
  - Matrix Multiplies:
    - M, N, K sizes should be multiples of 8
    - Larger K sizes make multiplications more efficient (amortize the write overhead)
    - Makes wider recurrent cells more practical (K is input layer width)
- **If you're designing models**
  - Make sure to choose layer widths that are multiples of 8
  - Pad input/output dictionaries to multiples of 8
    - Speeds up embedding/projection operations
- **If you're developing new cells**
  - Concatenate cell matrix ops into a single call

<https://youtu.be/i8-Jw48Cp8w>

# Download Docker Containers Framework compatible with TF32

Quantization  
TF32, FP16

Download Latest TensorFlow

<https://catalog.ngc.nvidia.com/orgs/nvidia/containers/tensorflow>

Download Latest PyTorch

<https://catalog.ngc.nvidia.com/orgs/nvidia/containers/pytorch>



NVIDIA NeMo  
Framework

# Thanks

## Any Question ?

